

Generalized Capsule Networks with Trainable Routing Procedure

Anonymous Authors¹

Abstract

Capsule Networks have been found to be effective in modeling spatial features with much fewer parameters and thus generalize better than CNNs. However, the routing procedure is designed to maximize the similarities between adjacent capsules rather than to maximize the global objective function. The routing number, which has to be set manually, is also difficult to optimize. We introduce the Generalized CapsNets (G-CapsNets) to overcome the above disadvantages by incorporating the routing procedure into the optimization process. G-CapsNets not only address the optimization of the coupling coefficients but also avoid the high computational overhead of the routing procedure. We implement two versions of G-CapsNets, fully-connected and convolutional on Caffe and compare them with CapsNets. The experiments show that G-CapsNets achieve better performance on MNIST & CIFAR10, and generalize better on GAN-generated synthetic images. We also test G-CapsNets on robustness to white-box & black-box adversarial attack.

1. Introduction

Convolutional neural networks (CNN) (Lecun et al., 1998) have been found effective in many computer vision problems. CNNs work well because the initial layers can learn simple features whereas later layers are supposed to combine simple features into complex ones. However, a fundamental issue of CNNs is that the spatial hierarchies between simple objects and complex objects are not well captured, which can prevent CNNs from generalizing well. For example, to detect a face, a CNN will report a high confidence once it finds the components (eyes, nose, ears, etc.) of a face even if the relative spatial relationship between these components are not sensible (since the spatial relationships

do not contribute to the loss). It is still possible for CNNs to capture spatial relationships, but they need either deeper networks or many negative samples. For the same reason, CNNs do not generalize well across different angles of the same object, as found in (Hinton et al., 2018). CapsNets, on the other hand, are composed of capsules, which are groups of neurons whose activity vectors represent the instantiation parameters of a specific type of entity such as an object or an object part (Sabour et al., 2017). Capsules help naturally preserve the spatial hierarchies between object parts, and thus can achieve the same or better performance using many fewer parameters. The evidence of this can be found in the experiments of Sabour et al. (2017) in MNIST (LeCun & Cortes (1998)), CIFAR10 (Krizhevsky et al.) and smallNORB (LeCun et al. (2004)).

Given these proceeding properties of CapsNets, many researchers have explored various applications of CapsNets. For example, (Rawlinson et al., 2018) adapt CapsNets to unsupervised learning by sparsifying the last capsule layer, while (Jaiswal et al., 2018) treat CapsNets as the backbone for GANs and achieve lower error rate on both MNIST and CIFAR10. (O’Neill, 2018) replace normal neural networks with CapsNets for face verification. Similarly, (Duarte et al., 2018) and (Liu et al., 2018) also use CapsNets as backbones for video classification and object localization, respectively.

In spite of many successful applications, the routing procedure of CapsNets is computationally expensive, which limits their scalability. The issue becomes more serious when a CapsNet becomes deeper. For example, in (Sabour et al., 2017), the routing number has to be set as 3. A routing number that is smaller than 3 or larger than 3 would cause degradation of performance. For a 10-layer CapsNet, assuming we have to try 3 routing numbers for each layer, then we have to try 3^{10} combinations to find the best routing number assignment. This problem may significantly limit the scalability and efficiency of CapsNets in practice.

To overcome this issue, we propose Generalized CapsNets. By “generalized,” we mean we can train a CapsNet just like training a standard neural network, and the (local) optimality of the coupling coefficients is guaranteed. The key idea of CapsNet is to incorporate the routing procedure into the overall optimization procedure, which makes the coupling coefficients trainable instead of being calculated by dynamic

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the ICML 2019 Workshop “Understanding and Improving Generalization in Deep Learning”. Do not distribute.

	CapsNets	G-CapsNets
Optimization	SGD + heuristic	SGD
Routing procedure	agreement	training
Times of Routing	meta parameter	1
Coefficients sum	1	any number

Table 1. CapsNets versus G-CapsNets. For the top two lines, CapsNets adopt a routing procedure called “routing by agreement (dynamic routing, EM routing),” which is a heuristic algorithm. Thus the optimization of CapsNets is a hybrid of Stochastic Gradient Descent (SGD) + heuristic algorithm while the optimization of G-CapsNets is a pure SGD. For the bottom two lines, CapsNets require several times of the routing procedure to find appropriate coupling coefficients whose sum has to be 1 for each layer. On the contrary, the coupling coefficients in G-CapsNets are treated as normal parameters and thus can be any number and only need to be calculated once for each pass.

routing (Sabour et al. (2017)) or EM routing (Hinton et al. (2018)). The differences between CapsNets and G-CapsNet are shown in Table 1.

We also test the scalability, generalization, and robustness of G-CapsNets. We find that, compared to vanilla CNNs, CapsNets show better scalability and generalization by using fewer parameters. However, we also find that G-CapsNets is as vulnerable as vanilla CNNs to a white or black box adversarial attack according to our experiment. Similar results can be found in (Marchisio et al., 2019; Yoon, 2017). This is different from (Hinton et al., 2018) who found that CapsNets are more robust to white-box attacks. We conjecture that the robustness of (Hinton et al., 2018) comes from introduced randomness during the routing procedure rather than the network itself.

2. Related work

Some researchers also found that the routing procedure is computationally expensive and try to improve it. For example, (Wang & Liu, 2018) introduce a regularizer based on KL divergence to minimize the clustering loss between capsules of adjacent layers. (Li et al., 2018) try to soften the issue of computational cost by adopting two extra branches to approximate the routing process. Above work’s contribution is proposing an alternative procedure to replace the expensive “routing by agreement”. However, the approximation of the routing procedure is restricted to calculating the coupling coefficients accurately which is not necessarily consistent with the final objective function. G-CapsNets, on the other hand, simplify the routing procedure and make it serve the whole objective function directly.




Figure 1. The structure of G-CapsNets.

3. Our approach: Generalized CapsNet

3.1. Structure of Generalized CapsNet

Similar to CapsNets, each capsule layer of G-CapsNets has two operations: capsule transformation and capsule routing. As Figure 1 shows, capsule transformation performs dimension transformation between adjacent capsule layers, while capsule routing combines the transformed capsules.

3.1.1. CAPSULE TRANSFORMATION & ROUTING

Capsule transformation happens between adjacent capsule layers, converting one type of capsule \mathbf{u}_i into another type of capsule (\mathbf{u}_{ji}) through a matrix transformation operation (\mathbf{W}_{ij}). Namely, $\mathbf{u}_{ji} = \mathbf{W}_{ij}\mathbf{u}_i$. In theory, we can transform any type of capsule into any other type of capsule. The capsules can be tensors of any shape as long as we have the appropriate transformation matrix. For example, in Sabour et al. (2017), 8-dimension capsules are transformed into 16-dimension capsules, while in Hinton et al. (2018), 4×4 capsules are transformed into 4×4 capsules.

Capsule routing ensures capsules in lower layers are scaled and sent to their parent capsules in higher layers. From another point of view, capsule routing can be considered as a clustering procedure in which the capsules in the upper layer are centers while the capsules in the lower layer are points that need to be chosen. Capsule routing combines information to forge new capsules $\mathbf{v}_j = \sum_i c_{ij}\mathbf{u}_{ji}$.

Note that the coupling coefficients c_{ij} are not acquired by the “routing-by-agreement” but depend on the optimization target. Thus the coupling coefficients could be any number (depending on the loss term and the regularizer term of the whole network). In other words, G-CapsNets choose whatever routing coefficients best fit the loss function, and thus the sum of the coupling coefficients ($\sum_i c_{ij}$) do not have to be 1.

Although the two operations of G-CapsNets are similar to CapsNets, there exists a fundamental difference between them. The routing procedure of CapsNets serves as the opti-

mization of similarities between capsules while the routing procedure in G-CapsNets serves as a step for the global optimization target.

3.2. Activation & loss function

The idea of a squash function is to map the length of a capsule to a number between 0 and 1. We adopt $\hat{v}_j = \left(1 - \frac{1}{e^{\|\mathbf{v}_j\|}}\right) \frac{\mathbf{v}_j}{\|\mathbf{v}_j\|}$ as suggested by Edgar et al. (2017), rather than the one in (Sabour et al., 2017) to achieve faster convergence. The margin loss function (see Equation 1) is the same as (Sabour et al., 2017).

$$L_k = T_k * \max(0, m^+ - \|\mathbf{v}_k\|)^2 + \lambda * (1 - T_k) * \max(0, \|\mathbf{v}_k\| - m^-)^2. \quad (1)$$

4. Experiments

4.1. G-CapsNets on MNIST & CIFAR10

4.1.1. FULLY CONNECTED G-CAPSNETS ON MNIST

We adopt the same baseline as described in Sabour et al. (2017) which has three convolutional layers of 256, 256, and 128 channels. The kernels and strides are 5x5 and 1. The last convolutional layer is followed by two fully-connected layers of size 328 and 192. The final layer is a 10-class softmax classifier.

For G-CapsNets, we also adopt the same architecture as in Sabour et al. (2017). The first convolutional layer outputs 256 feature maps. The second convolutional layer outputs 256 feature maps or $32 \times 6 \times 6$ 8D capsules. For the final layer, we replace the dynamic routing procedure with our trainable routing procedure. Please check our released code or the original paper written by Sabour et al. (2017) for more details.

We call the capsule structure here (and the one in Sabour et al. (2017)) fully-connected CapsNet since each capsule in the higher layer connects to every capsule in the lower layer. As Table 2 shows, no matter the reconstruction involved, G-CapsNets can always achieve better performance than vanilla CNNs as well as CapsNets (Sabour et al., 2017). Note that the accuracy of the baseline and the CapsNets reported here is lower than in Sabour et al. (2017) due to differences in the data augmentation (pixel-shifting) adopted by Sabour et al. (2017) as well as different frameworks (Caffe versus TensorFlow).

4.1.2. CONVOLUTIONAL G-CAPSNETS ON MNIST

Similar to the structure in Hinton et al. (2018), we build a convolutional version of G-CapsNets. The same type of capsules of different positions share the same transforma-

Algorithm	error rate(%)	param #
vanilla CNN	$0.91 \pm 0.09 / 36.54 \pm 0.78$	35.4M/9.6M
CapsNets**	$1.24 \pm 0.10 / 40.88 \pm 0.67$	8.2M/2.67M
FC G-CapsNets	$0.67 \pm 0.05 / 32.70 \pm 0.47$	8.2M/2.67M
FC G-CapsNets*	$0.66 \pm 0.03 / -$	6.8M/-
Conv G-CapsNets	$0.84 \pm 0.09 / 33.53 \pm 0.52$	6.9M/2.67M
Conv G-CapsNets*	$0.86 \pm 0.13 / -$	5.5M/-
Multi G-CapsNets	$- / 34.29 \pm 0.33$	- / 716K

Table 2. Error rate versus number of parameters on MNIST. Note that “*” means no reconstruction, and “CapsNets**” is the algorithm that comes from (Sabour et al., 2017).

tion matrices. We use a 6x6 kernel for the last capsule layer and a 4x4 “matrix” to transform capsules (from 4x4 matrix capsules to 4x4 matrix capsules). As Table 2 shows, convolutional G-CapsNets achieve better performance compared to the baseline by using fewer parameters.

4.1.3. G-CAPSNETS ON CIFAR10

For CIFAR10, we build a similar G-CapsNet structure as the one used for MNIST. There are two differences. (1) we adopt 64 feature maps (other than 256 feature maps) in the first convolutional layer, and (2) we adopt 8x8 kernel for the capsule layer and transform each capsule by using an 8x8 matrix. The baseline shares the same structure as the G-CapsNets used here for the first two layers, and has two fully-connected layers with an output of 512 and 10 separately. As Table 2 shows, both fully-connected CapsNets, and convolutional CapsNets achieve better performance than the baseline as well as CapsNets. We also develop a multi-layer version of G-CapsNets, which also achieves better performance than both vanilla CNN and CapsNets, as Table 2 shows. Please refer to the supplementary appendix for more details.

4.2. The generalization of G-CapsNets

CapsNets are better at capturing the relationship of different spatial features, so it makes sense to believe that CapsNets have a better generalization ability. To verify if this is true, we use AC-GAN (Auxiliary Classifier Generative Adversarial Network) proposed by (Odena et al., 2017) to generate synthetic images. Specifically, we generate 4000 artificial MNIST-like images (400 images for each digit), as Figure 2 shows. AC-GAN encodes both class labels as well as noise to generate synthetic images. The discriminator of AC-GAN also outputs a distribution over the source and the class labels. Please refer to the original paper for details. The discriminator contains four convolutional layers which have 32, 64, 128, and 256 feature maps. Each convolutional layer is followed by a leaky ReLU layer and a dropout layer.

The generator first uses a fully-connected layer to map the latent vector (100, 1) to a (3, 3, 384) tensor, then up-samples the tensor to a (7, 7, 192) tensor, a (14, 14, 96) tensor and a (28, 28, 1) tensor. The learning rate (2e-4) and beta1 (0.5) are those recommended by Radford et al. (2015).

These images are recognizable for humans but not the same as the original images in MNIST (in terms of distribution) since they are generated in the early stage of AC-GAN. Being recognizable means these images have enough clues to tell their classes and being generated in the early stage means they do not share the same distribution as the authentic MNIST digits. Both properties make these images perfect for testing generalization.

The underlying assumption of G-CapsNets is that they are better at capturing the spatial features in an image and thus they need far fewer parameters than standard neural networks. As Table 3 shows, both the fully connected version and convolution version of G-CapsNets achieve better performance than the baseline, which uses more parameters. However, both the baseline and the G-CapsNets show no significant difference when the number of epochs is larger than 3. We argue that the reason is that the generated images have already converged to the same distribution as in MNIST. Note that it makes more sense to compare the performance over three network structures rather than across epochs since the generated images vary each time.

Epochs	Baseline	FC G-CapsNets	Conv G-CapsNets
1	19.80	23.50	19.725
2	5.60	3.25	4.05
3	0.725	0.00	0.70
4	0.0005	0.00	0.00
5	0.00275	0.00	0.001

Table 3. Error rate on MNIST-like images generated by AC-GAN. FC CapsNet is full connected G-CapsNets and Conv G-CapsNets refers to convolutional G-CapsNets. The number of parameters of the baseline, FC G-CapsNets, and Conv G-CapsNets are 35.4M, 6.8M, and 5.5M respectively.

4.3. The robustness of G-CapsNets on black box attack

Hinton et al. (2018) claimed that CapsNets are comparable to CNNs regarding robustness to black-box attacks despite using fewer parameters. Our experiments also support this claim. Specifically, we adopt LeNet as the substitute model to generate perturbations for each testing images (10K testing images in MNIST) based on FSGM (Goodfellow et al. (2015)). We restrict the maximum perturbation for each pixel to be 8 ($L_\infty \leq 8$). We found that the accuracy of the baseline, the fully-connected G-CapsNets, and the convolutional G-CapsNets drops sharply (to 11.35%, 8.92%, 11.35% after attacking). We thus conclude that G-CapsNets




Figure 2. The generated MNIST-like images. From top to bottom, the number of training epochs is 1, 2, 3, 4 and 5.

is as vulnerable as standard neural networks.

4.4. The robustness of G-CapsNets on white-box attack

Hinton et al. (2018) found that CapsNets are robust to white-box adversarial attack due to numerical instability (Nayebi & Ganguli (2017)) as well as the smaller percentage of zero values in the gradient. However, their testing was based on FGSM Goodfellow et al. (2015) which is not a strong attack technique. Inspired by Universal Adversarial Perturbation (UAP) (Moosavi Dezfooli et al. (2017)), we trained a generative model to generate universal perturbations during training (offline), and then applied the generated perturbations to all testing images. We found that G-CapsNets do not show stronger robustness to the attack. Please check the supplementary appendix for more details.

5. Conclusion

G-CapsNets incorporate the capsule routing procedure into the whole optimization process, avoids the need to set routing times for each capsule layer and the optimization is addressed. The two versions of G-CapsNets, fully-connected G-CapsNets and convolutional G-CapsNets achieve better performance on MNIST and CIFAR10 compared to both CNNs and CapsNets. G-CapsNets also show good scalability as well as generalization ability. Finally, we evaluated the robustness of G-CapsNets against both the white-box attack and black-box attack. For all the measurements we adopt in this paper, G-CapsNets achieve either better (concerning the accuracy, number of parameters, generalization ability, scalability) or comparable (regarding the white-box attack and black-box attack) performance.

References

- Duarte, K., Rawat, Y. S., and Shah, M. Videocapsule-net: A simplified network for action detection. *CoRR*, abs/1805.08162, 2018. URL <http://arxiv.org/abs/1805.08162>.

- Edgar, X., Selina, B., and Yang, J. Capsule network performance on complex data. *CoRR*, 2017. URL <https://arxiv.org/abs/1712.03480>.
- Goodfellow, I., Shlens, J., and Szegedy, C. Explaining and harnessing adversarial examples. In *International Conference on Learning Representations*, 2015. URL <http://arxiv.org/abs/1412.6572>.
- Hinton, G. E., Sabour, S., and Frosst, N. Matrix capsules with EM routing. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=HJWLfGWRb>.
- Jaiswal, A., AbdAlmageed, W., Wu, Y., and Natarajan, P. CapsuleGAN: Generative Adversarial Capsule Network. *CoRR*, February 2018.
- Krizhevsky, A., Nair, V., and Hinton, G. Cifar-10. URL <http://www.cs.toronto.edu/~kriz/cifar.html>.
- LeCun, Y. and Cortes, C. MNIST handwritten digit database. 1998. URL <http://yann.lecun.com/exdb/mnist/>.
- Lecun, Y., Bottou, L., Bengio, Y., and Haffner, P. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, Nov 1998. ISSN 0018-9219. doi: 10.1109/5.726791.
- LeCun, Y., Huang, F., and Bottou, L. Learning methods for generic object recognition with invariance to pose and lighting. In *IEEE Conference on Computer Vision and Pattern Recognition*, 2004.
- Li, H., Guo, X., Dai, B., Ouyang, W., and Wang, X. Neural network encapsulation. *ECCV*, 2018.
- Liu, W., Barsoum, E., and Owens, J. D. Object localization and motion transfer learning with capsules. *CoRR*, abs/1805.07706, 2018. URL <http://arxiv.org/abs/1805.07706>.
- Marchisio, A., Nanfa, G., Khalid, F., Abdullah Hanif, M., Martina, M., and Shafique, M. CapsAttacks: Robust and Imperceptible Adversarial Attacks on Capsule Networks. *CoRR*, art. arXiv:1901.09878, Jan 2019.
- Moosavi DeZfooli, S. M., Fawzi, A., Fawzi, O., and Frossard, P. Universal adversarial perturbations. *Proceedings of 2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 9, 2017. doi: 10.1109/Cvpr.2017.17. URL <http://infoscience.epfl.ch/record/226156>.
- Nayebi, A. and Ganguli, S. Biologically inspired protection of deep networks from adversarial attacks. *CoRR*, March 2017.
- O’Neill, J. Siamese Capsule Networks. *CoRR*, May 2018.
- Odena, A., Olah, C., and Shlens, J. Conditional image synthesis with auxiliary classifier GANs. In Precup, D. and Teh, Y. W. (eds.), *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pp. 2642–2651, International Convention Centre, Sydney, Australia, 06–11 Aug 2017. PMLR. URL <http://proceedings.mlr.press/v70/odena17a.html>.
- Radford, A., Metz, L., and Chintala, S. Unsupervised representation learning with deep convolutional generative adversarial networks. *CoRR*, abs/1511.06434, 2015. URL <http://arxiv.org/abs/1511.06434>.
- Rawlinson, D., Ahmed, A., and Kowadlo, G. Sparse unsupervised capsules generalize better. *CoRR*, abs/1804.06094, 2018. URL <http://arxiv.org/abs/1804.06094>.
- Sabour, S., Frosst, N., and Hinton, G. E. Dynamic routing between capsules. *CoRR*, abs/1710.09829, 2017. URL <http://arxiv.org/abs/1710.09829>.
- Wang, D. and Liu, Q. An optimization view on dynamic routing between capsules, 2018. URL <https://openreview.net/forum?id=HJjtFYJDf>.
- Yoon, J. Adversarial Attack to Capsule Networks, November 2017. URL https://github.com/jaesik817/adv_attack_capsnet.

Layer	Layer parameters	param #
Conv#1	(7, 7, 64), 1	9.4K
Conv#2	(7, 7, 128), 2	401.4K
Conv Caps Trans#1	(4, 4, 8), 1, (8, 8)	131.1K
Conv Caps Routing#1	(3, 3, 16), (7, 7, 8)	56.4K
Conv Caps Trans#2	(3, 3, 16), 2, (2, 2)	4.6K
Conv Caps Routing#2	(3, 3, 8), (3, 3, 16)	10.4K
FC Caps Trans#1	(16, 3, 3, 8), (8, 10)	92.2K
FC Caps Routing#1	(16, 3, 3, 8), 10	11.5K
Total	-	716K

Table 1. The structure of Multi-layer G-CapsNets.

6. Supplementary Material

6.1. The structure of Multi-layer G-CapsNets on CIFAR10

To test the scalability of G-CapsNets, we build a network with two convolutional layers, two convolutional capsule layers, and one full connected layer, as Table 1 shows. We adopt one ReLU layer and one squash layer after each convolutional capsule layer and full connected capsule layer. Take the ‘Conv Caps Transform#1’ as an example, (4, 4, 8) means the kernel size is (4, 4) and the number of capsule feature maps is 8. The following number is the stride. The last (8, 8) refers to the transformation matrix. Namely, the G-CapsNets transform capsules (1x8) on the lower layer to capsules (1x8) on the upper layer. ‘Conv Caps Routing#1’ follows the ‘Conv Caps Transform#1’, whose responsibility is combining the information from each tensor with the shape of (3, 3, 16), and we have a total number of 7x7x8 combinations. For the full connected CapsNet, take the ‘FC Caps Transform#1’ layer as example, the multi-layer G-CapsNets transform each tensor with a shape of (16, 3, 3, 8) on the lower layer to a new tensor with a shape of (16, 3, 3, 8) on the upper layer. Then the layer ‘FC Caps Routing#1’ combines each tensor of shape (16, 3, 3, 8) to a new tensor of shape (1, 8). Since CIFAR has ten classes, the output of the final layer has a shape of (10, 8).

6.2. White adversarial attack on G-CapsNet

The structure of the generative model is similar to a GAN but with the discriminator unchanged. Specifically, the input of the generator is a 100-dimension latent vector whose values are between 0 and 1. The latent vector layer is followed by three deconvolutional layers (note that we apply Batch Normalization after each de-convolutional layer). The final layer of the generative branch outputs a tensor with the same dimension as the input image. The discriminator is the classification model we need to test. The loss function is to minimize the difference between the logits of a clean image




Figure 1. The decreasing curve of three network structures (standard: standard neural network, fc: fully-connected G-CapsNets, conv: convolutional G-CapsNets) with white-box attack

and the logits of its manipulated version. The whole attack is untargeted, and we assign each under-attacked image a random incorrect label during training.

We apply this attack technique on the test set of MNIST which contains 1000 images. As Figure 1 shows, all three networks’ accuracy drop sharply after 100 attacking iterations. This result is not consistent with what (Hinton et al. (2018)) found with the weaker FGSM attack. These results suggest that CapsNets and CNNs are both vulnerable to strong white-box attacks like UAP.